

		Audio, production scale
<u>Qdrant</u>	Excellent, very flexible	Text + Image + Metadata hybrid search
Weaviate <u>cloud</u>	Good with Transformers + <u>hybrid</u>	RAG, text + image + metadata 'combo' } ←
Pinecone <u>local</u> <u>cloud</u>	Text-only focus (no native image)	Great for text-based RAG, not for images
Vespa.ai	Enterprise-heavy multimodal <u>infra</u>	Multimodal + advanced query logic
AWS OpenSearch (KNN)	{ — — — }	{ }
Azure AI Search		
GCP Vertex AI Vector Search		

The **difference is not in creating the vector**; the difference lies in the **features and ecosystem of the vector database**.

Pinecone

```
{"id": "doc-1", "values": [0.22, 0.88, ...]}
```

Weaviate

BS4

```
{  
  "class": "Product",  
  "properties": {  
    "image": "path.jpg",  
    "text": "Best shirt in market"  
  }  
}
```

embedding

Multi modal data

Vector DB	Verdict	Use-case Fit
Milvus (Zilliz)	good for multimodal	Image + Text + Video + Audio, production scale
<u>Qdrant</u>	Excellent, very flexible	Text + Image + Metadata

YouTube caption search

Video + Audio + Text

Resume screening from PDF + profile

Text + Table + Image (logos etc.)

Search from screenshots

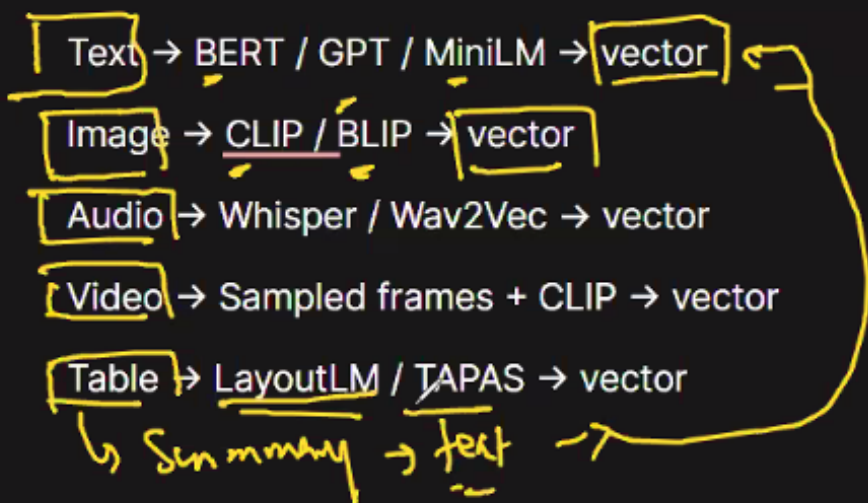
Image → OCR → Text → Semantic Search

Voice-commanded document search

Audio → Text + Query understanding

Medical Report Q&A

Document (PDF) = Text + Charts + Tables



The **difference is not in creating the vector**; the difference lies in the **features and ecosystem of the vector database**.

Video → moving visuals with audio

Documents → combination of text + images

Use-case

Modalities Involved

ChatGPT Vision

Image + Text

YouTube caption search

Video + Audio + Text

Resume screening from PDF + profile

Text + Table + Image (logos etc.)

Search from screenshots

Image → OCR → Text → Semantic Search

Voice-commanded document search

Audio → Text + Query understanding

Medical Report Q&A


Document (PDF) = Text + Charts + Tables

Text → BERT / GPT / MiniLM → vector


Image → CLIP / BLIP → vector

Modalities = Different types of input data
Text → plain sentences, paragraphs
Image → photos, diagrams, charts } <u>image</u>
Audio → voice, speech, music = <u>Audio</u>
Video → moving visuals with audio ← <u>Frame + Audio</u>
Documents → combination of text + images ← <u>PI</u>

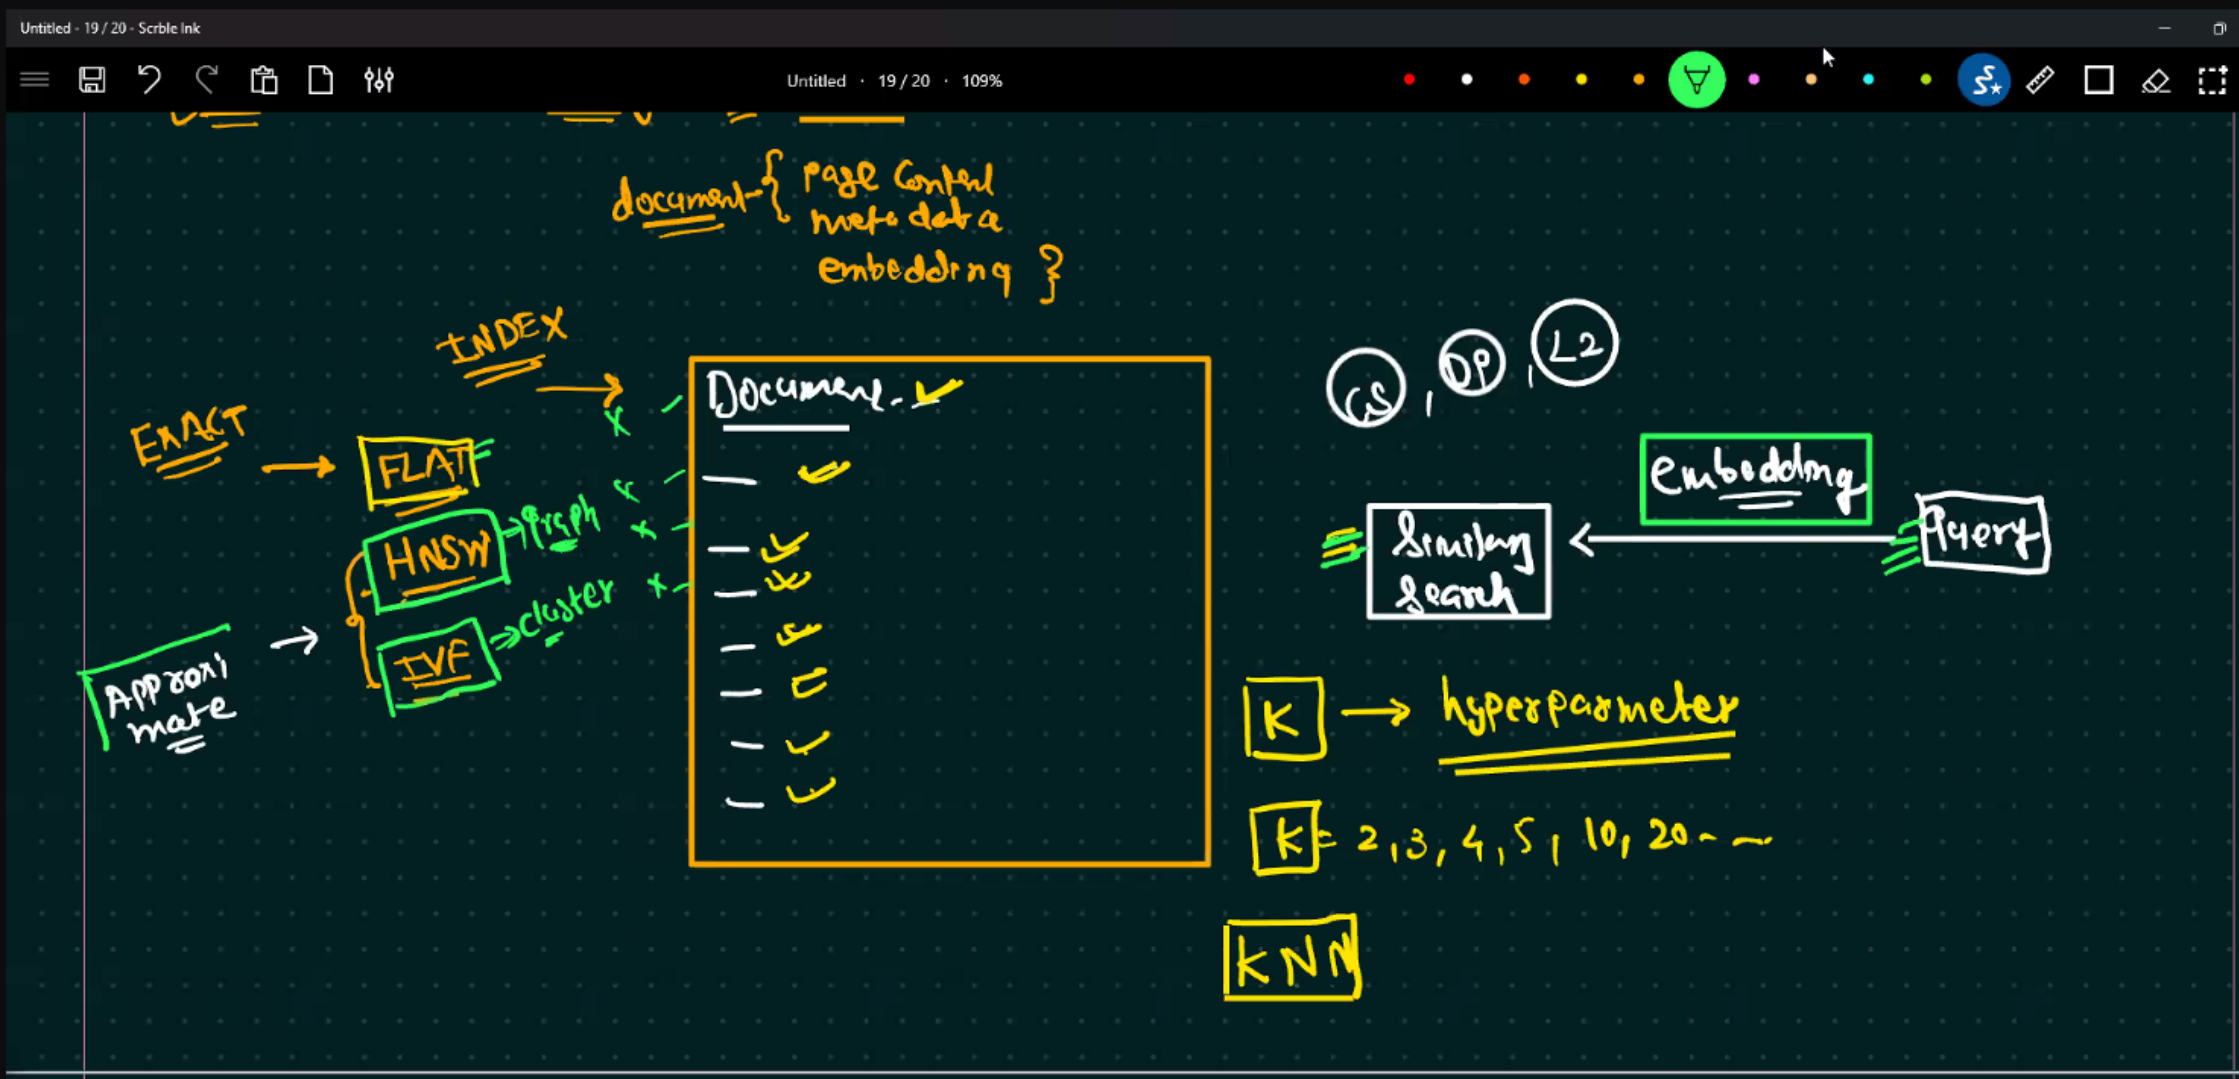
Use-case	Modalities Involved
ChatGPT Vision	Image + Text
YouTube caption search	Video + Audio + Text
Resume screening from PDF + profile	Text + Table + Image (logos etc.)



= KNN



= ANN



Vector Database ...

Document

Both

Canvas

Ctrl K

Share



Vector Database

(LangChain etc.)

Use-case fit

Small, fast-only

Dev/PoC

Production scale

Vendor

Local Support

Cloud Hosted

FAISS

Yes

No

Quadrant

Yes

Yes

Weaviate

Yes

Yes

Pinecone

No (API only)

Only

Milvus

Yes

(Zilliz)

Chroma

Yes

No

[Langchain Vector DB Comparison](#)

[VectorDB Comparison](#)

Term

Full Form

Meaning in Simple Terms

+ - | T ~ :≡ ≡≡ ≡≡ | </> ☞ ≡ neighbors

Exact search

: Return top K



Upgrade to unlock unlimited files, more AI credits, and custom icons settings & [more](#).

Join Slack Community

Upgrade

Vector Database

Document

Both

Canvas

Ctrl K

Share

Ctrl K

Ctrl K

Vector Database

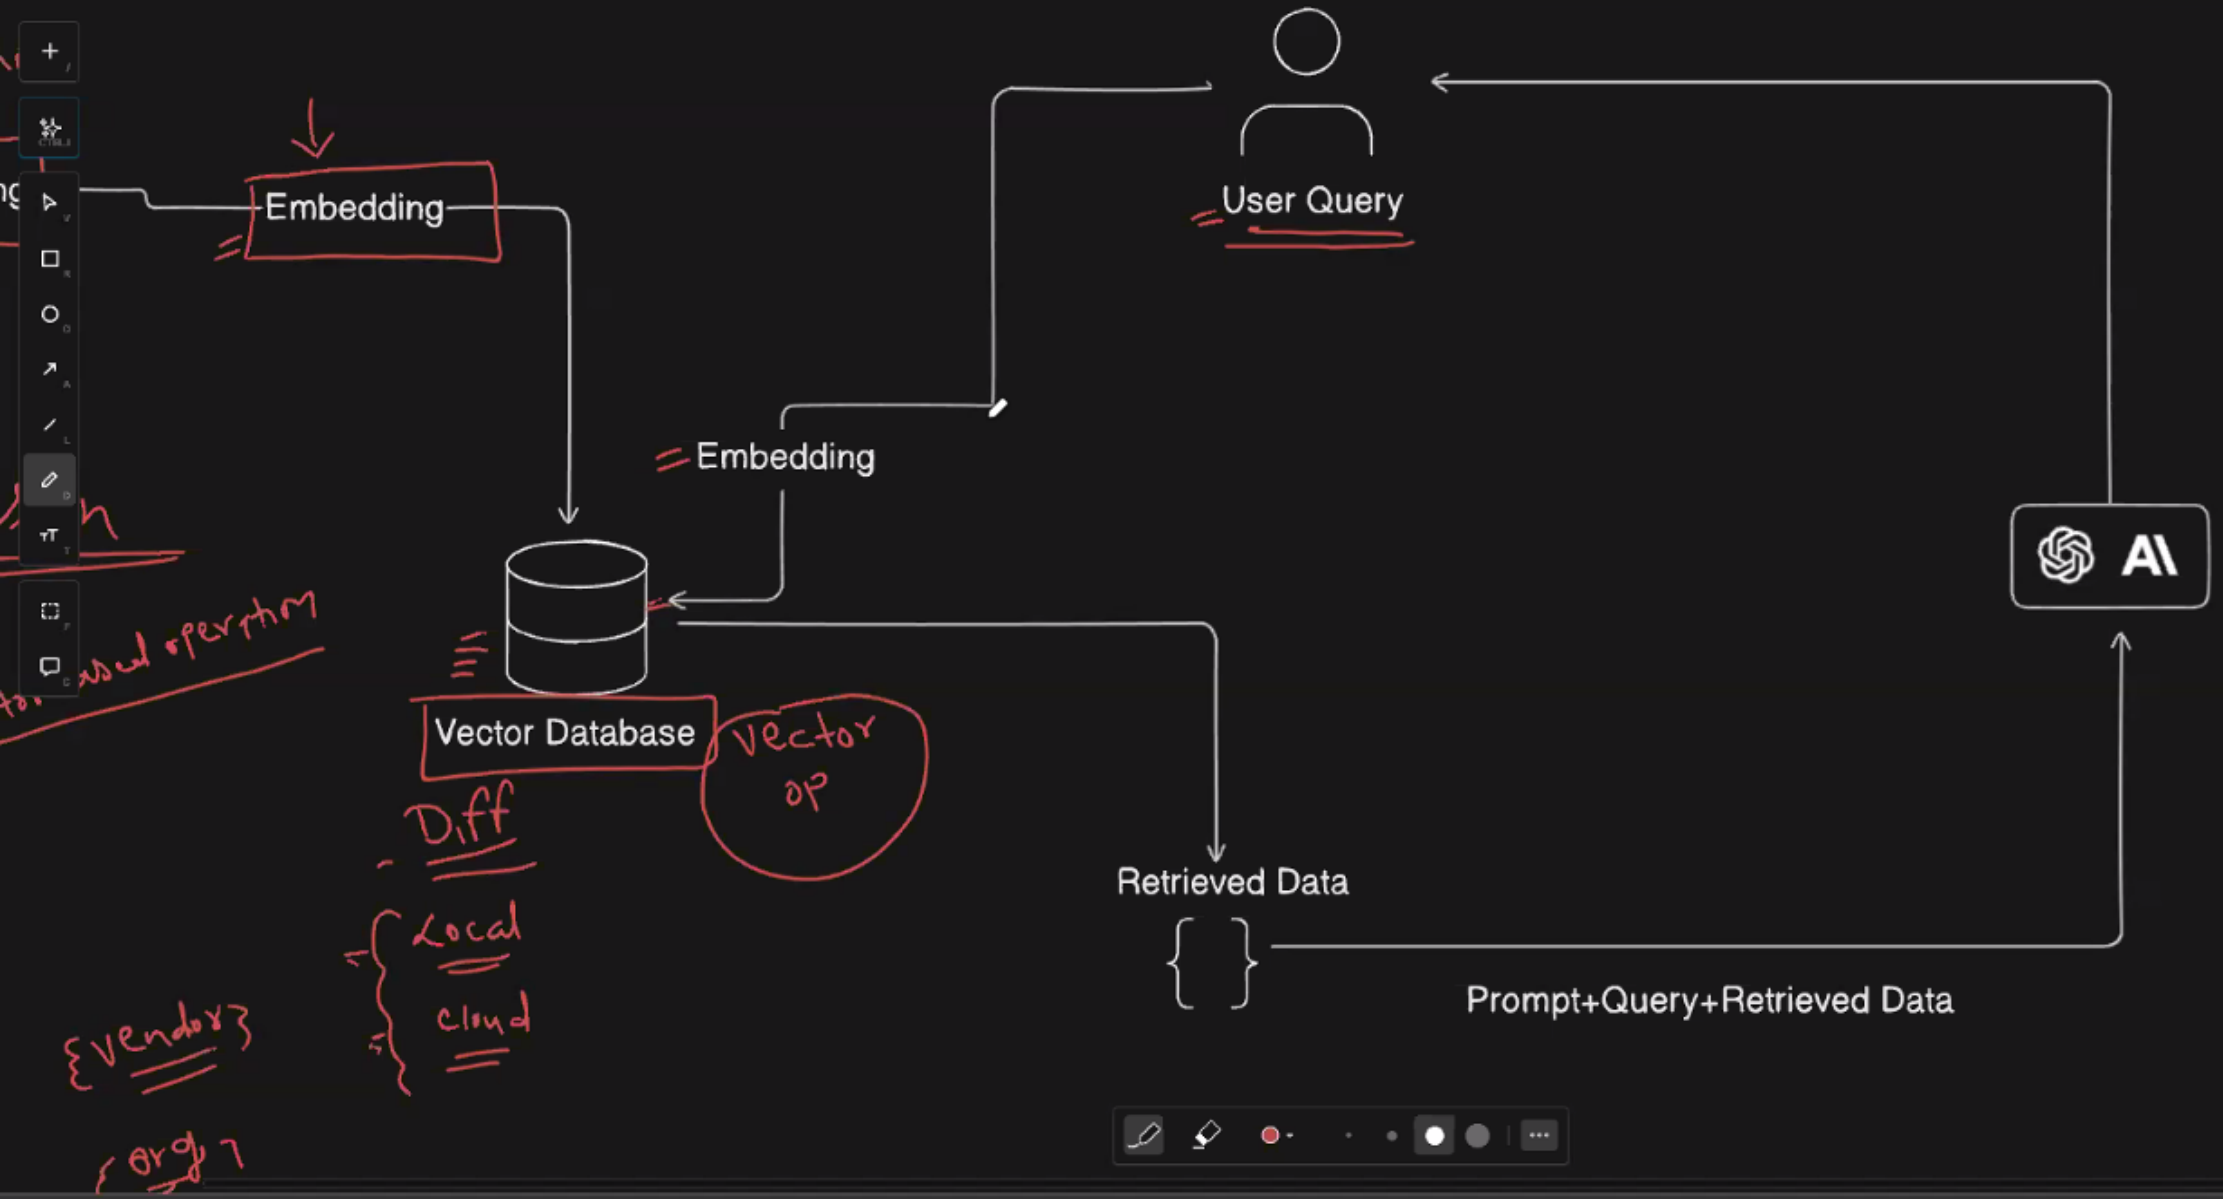
Type	Description	Example DBs
Type	Description	Example Databases
Local	Runs on your own laptop or server, either on the file system or via Docker	FAISS, Qdrant (local), Weaviate
In-Memory	Stores data only in RAM - data is lost when the system restarts	FAISS, Qdrant (local), Weaviate
Cloud Hosted	Fully hosted by vendors - supports persistence, scaling, APIs, monitoring, etc.	Pinecone, Qdrant Cloud, Weaviate Cloud, Milvus (Zilliz Cloud)

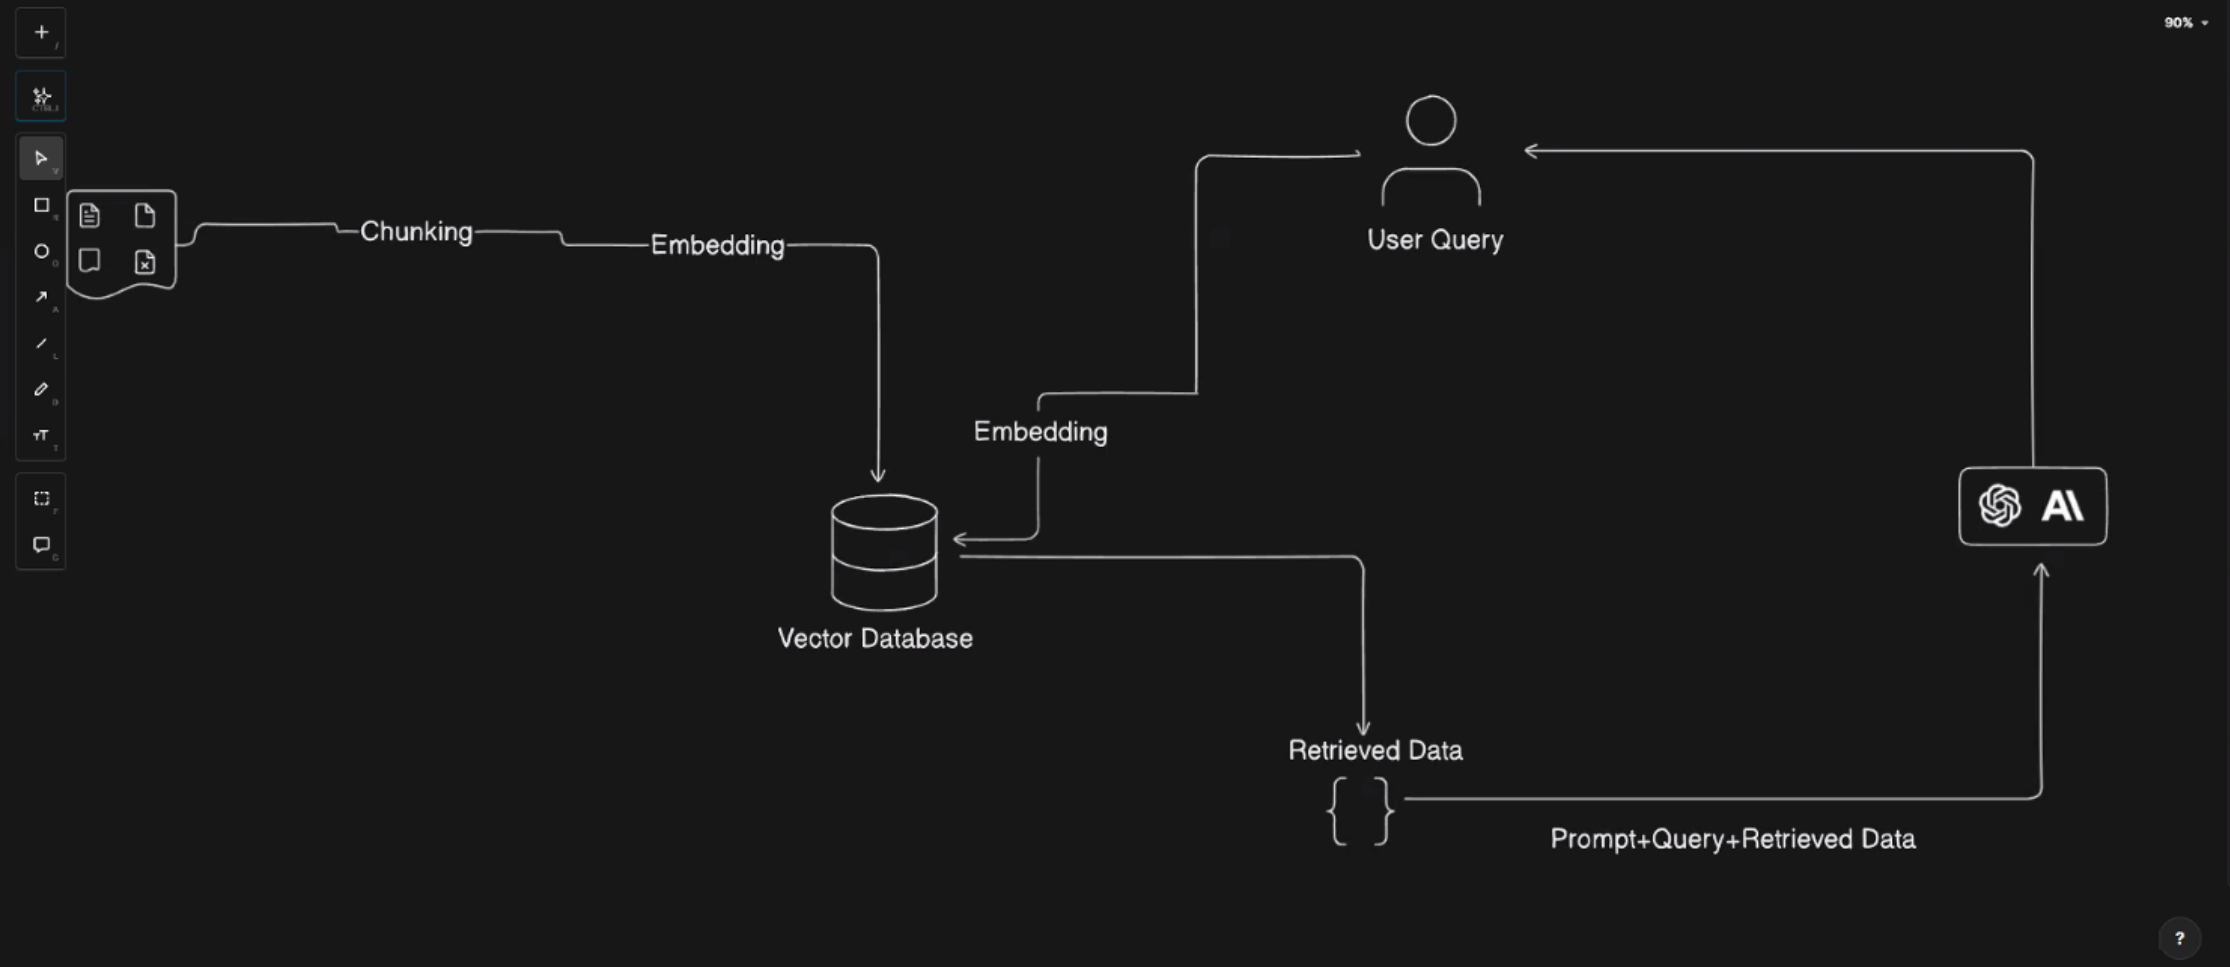
Local: useful for **experiments**, dev, prototypes, or small deployments

Upgrade to unlock unlimited files, more AI credits, and custom icons settings & more.

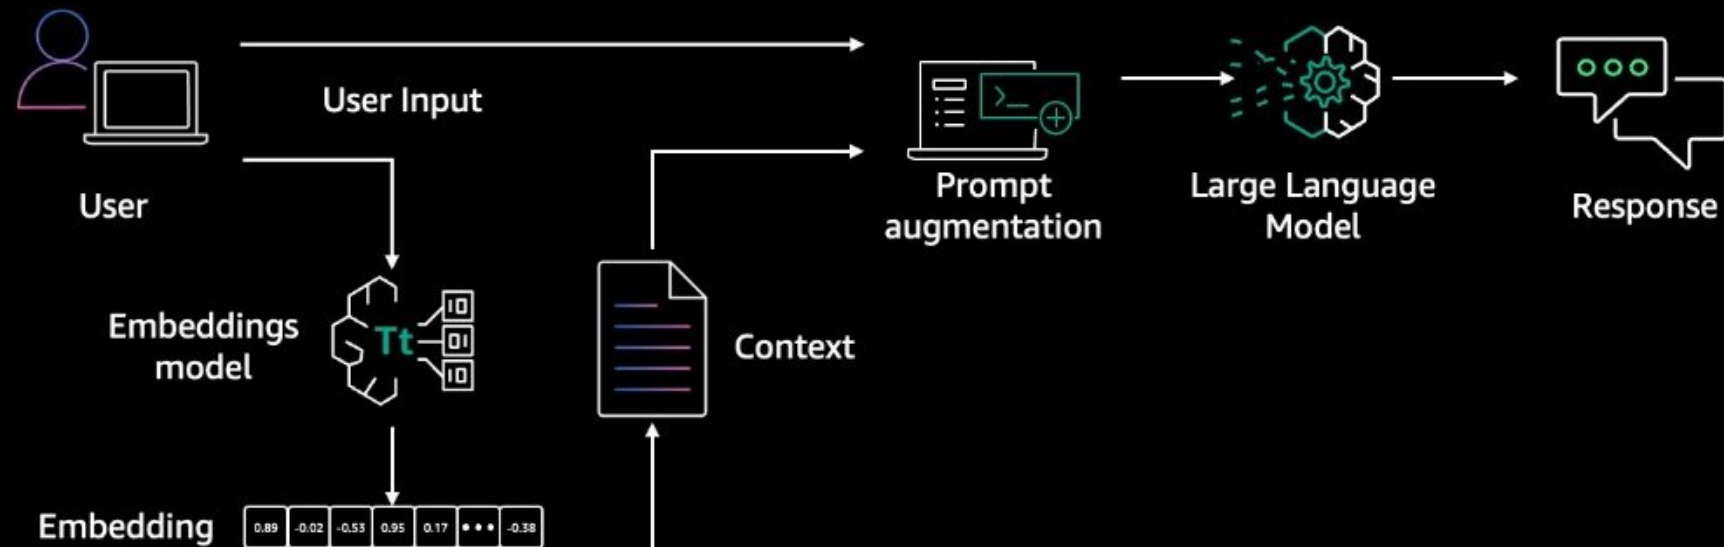
Vector Database

Document Both Canvas






Text Generation Workflow



Data Ingestion Workflow



Upgrade to unlock unlimited files, more AI credits, and custom icons settings & [more](#).

Vector Database ...

Document

Both

Canvas

- Real-time updates
- Monitoring & backups



Vector Database

Feature	In-Memory	Local Persistent	Cloud Hosted
<u>Data survives restart</u>	No ✗	Yes ✓	Yes ✓
<u>Scales across nodes</u>	No ✓	Limited	Yes
Team access (multi-user)	No	Manual	Yes
Easy integration (LangChain etc.)	SDKs	SDKs	SDKs/APIs ready
Use-case fit	Small, fast-only	Dev/PoC	Production scale

Vendor	Local Support	Cloud Hosted
FAISS	Yes	No
Quadrant	Yes	Yes